



Web Stream Processing with OntopStream

The Web Conference 2022


Pieter Bonte
Matteo Belcao
Marco Balduini
Emanuele Della Valle

OntopStream

Ontology-Based Data Access

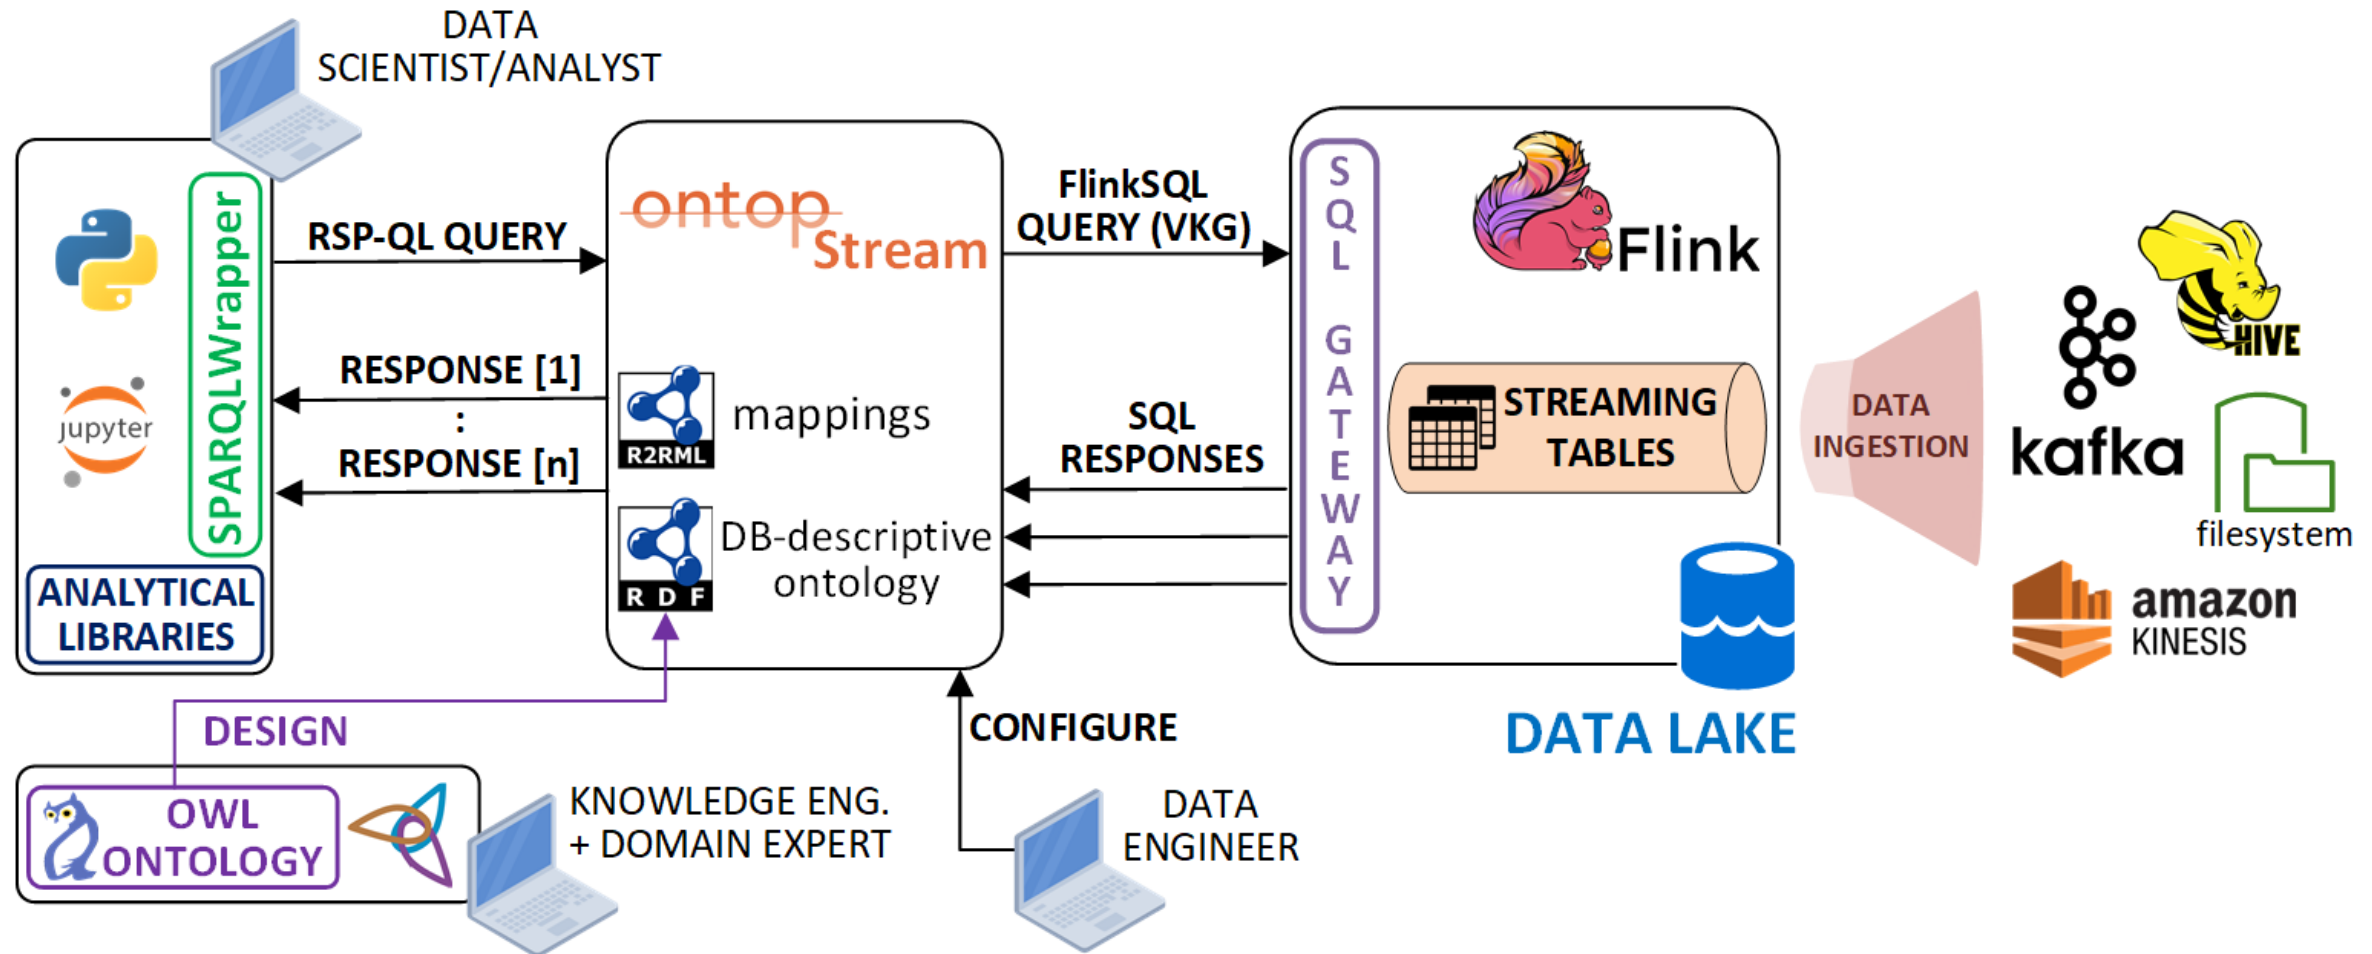
- **Ontology-Based Data Access (OBDA)** softwares aim to solve data integration problems...
- **Virtual Knowledge Graph (VKG)** approach:
 - additional semantic layer on top of the data
 - relational data sources abstraction, exposed as RDF triples
 - SPARQL queries to access the data
 - automatic SPARQL → SQL query rewritings

KG-Empowered Continuous Analytics



KG-Empowered Continuous Analytics

Streaming-VKGs as a bridge between Stream Processing and Semantic Techs




OntopStream

- Developed as an extension of the Ontop OBDA system (Java)
- Query relational data streams
 - stored and managed in Apache Flink dynamic tables
 - with RSP-QL continuous queries (windowed / not windowed)
- Get RDF streams of responses
- Two distributions:
 - OntopStream-CLI
 - OntopStream-Endpoint (only HTTP calls)

OntopStream: design decisions

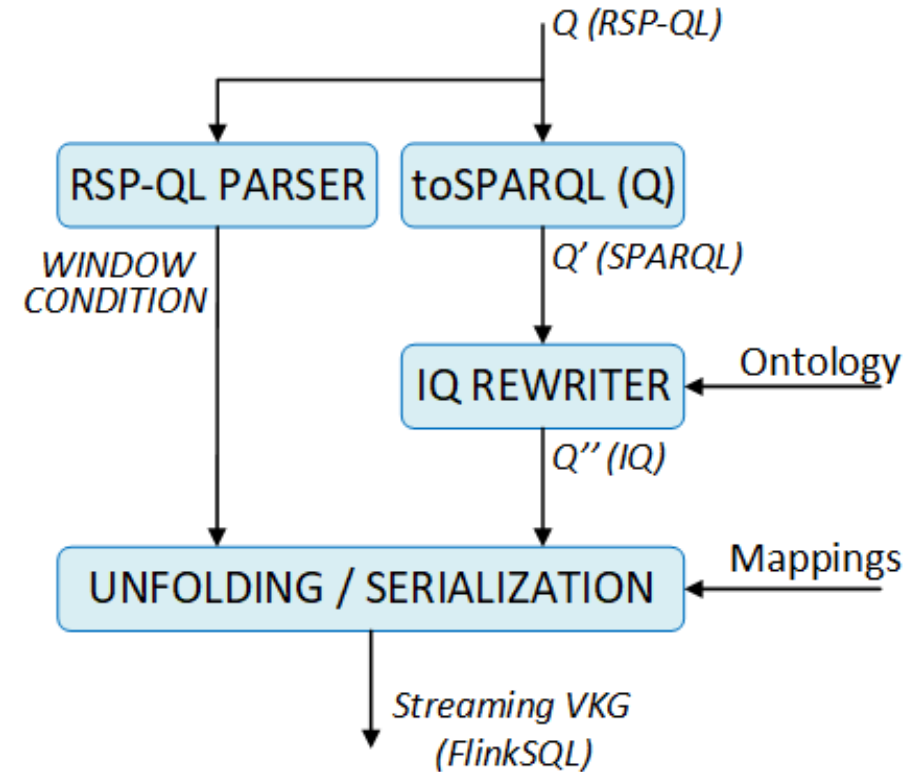
- paradigm shift from traditional OBDA to Streaming-OBDA
- design decisions:
 1. extend the **Flink JDBC driver**
 2. re-design part of the **ontop-engine** to accept **RSP-QL** queries
 3. Streaming Virtual Knowledge Graph query rewriting approach
 4. include support for **RDF streams** of query outputs

Streaming Virtual Knowledge Graphs in OntopStream



Streaming Virtual Knowledge Graph query rewriter

- **rsp4j** parser to extract window conditions
- Intermediate Query rewriter unchanged
- **IQ** representation:
 - created w.r.t to the Ontology **O**
 - unfolded in a **Streaming VKG** tree
- Each tree node corresponds to a pseudo-SQL statement
- Streaming VKG serialization in a **FlinkSQL** query, add window condition **W** if existing



Tutorial: pipeline setup

Starting-up the resources

- Requirements: *docker* and *docker-compose* ([installation guide](#))
- Clone the repository: `git clone https://github.com/pbonte/WSP-TheWebConf2022Tutorial.git`
- Start the tutorial environment
 - Open the hands-on folder: `cd exercises/part3`
 - Streaming resources (Flink, Kafka) and JupyterLab:
Windows: `docker-compose -f flink-kafka-win.yml up -d`
UNIX/Mac: `docker-compose -f flink-kafka-unix.yml up -d`
 - Flink JDBC Gateway:
`docker exec -it sql-client /opt/flink-sql-gateway-0.2-SNAPSHOT/bin/sql-gateway.sh --library /opt/sql-client/lib`
Note: keep the JDBC endpoint alive until you need the service (don't close the terminal window)

Business Scenario: Rental Company

A car rental company has recently decided to **unify the information systems** of **two branches** using ontology-based data access techniques.

Both the branches:

- have a real-time data management infrastructure
- store the rental records in Kafka topics

However, they handles the data differently:

- Branch A uses two separate Kafka topics for trucks and cars
- Branch B stores all the rentals in a single topic, but the users' data are kept in a sperate topic

Business Requirements

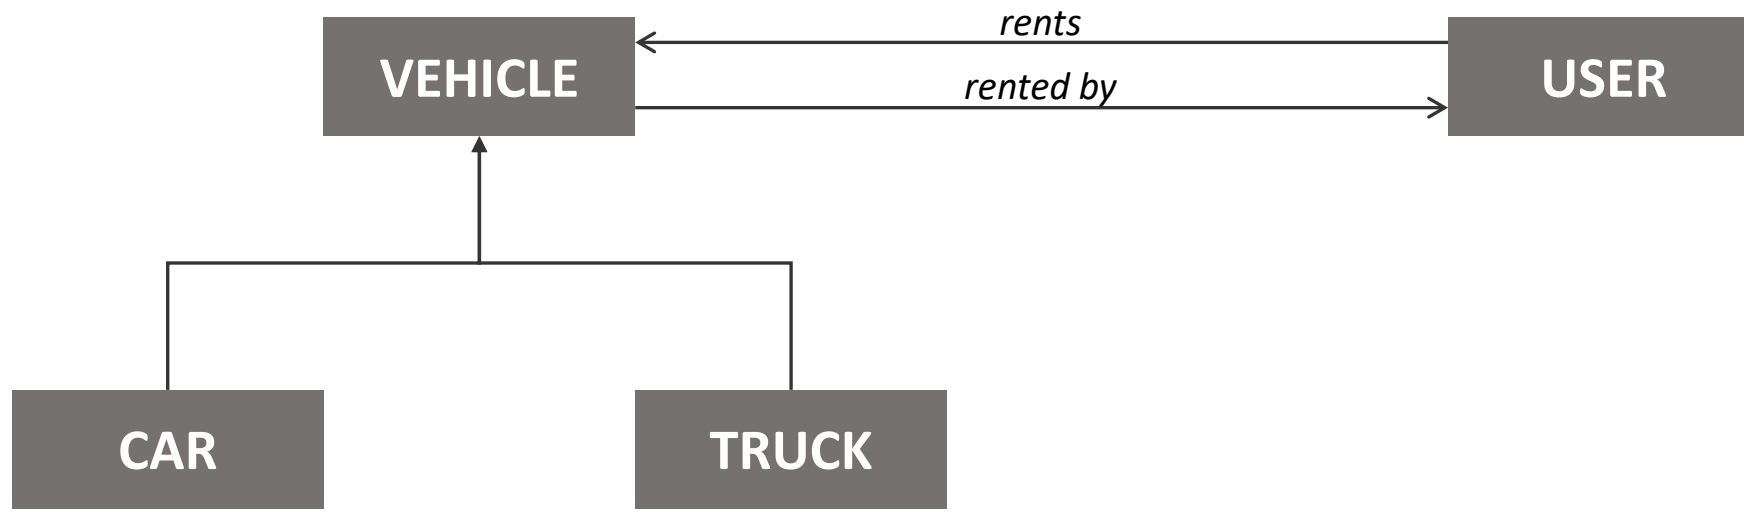
The company is booming, and has in plan to acquire soon new branches.

Therefore, the company wants to make the **integration process scalable**, so that can be easily extended to all its new branches

They need a data integration solution that:

- provides an **unified logical view** of their data
- enables to **query in real-time** their data
- can be used with **python notebooks** for further analyses

Kafka Topics: High-Level View



PROBLEM: the two branches uses different data structures...

Kafka topics: Branch A

user	rid	manufacturer	model	plate	status
Molly Davis	1	Fiat	Panda	FJ7PUJJ	START
Laura Baker	2	Tesla	Model S	JFGJ60A	START
William Diaz	3	Fiat	Tipo	FGL1X62	START
Molly Davis	1	Fiat	Panda	FJ7PUJJ	END
William Diaz	3	Fiat	Tipo	FGL1X62	END

DEALER1_CARS

user	rid	manufacturer	model	plate	status
Laura Baker	1	Iveco	Daily	HHST532	START
Wayne Flower	2	Fiat	Ducato	DM89JKD	START
Richard Tillman	5	Fiat	Ducato	JSDJFI3	START
Richard Tillman	5	Fiat	Ducato	JSDJFI3	END
Wayne Flower	2	Fiat	Ducato	DM89JKD	END

DEALER1_TRUCKS

Kafka topics: Branch B

DEALER2_VEHICLES

userID	rid	type	manufacturer	model	plate	status
3	1	Car	Audi	A3	DFU4HJF	START
4	2	Car	Mercedes	Classe C	784JD93	START
3	7	Truck	Mercedes	Vito	KD94KDS	START
3	1	Car	Audi	A3	DFU4HJF	END
6	8	Truck	Mercedes	Vito	012JKD0	START
3	7	Truck	Mercedes	Vito	KD94KDS	END

DEALER2_USERS

userID	name
1	Douglas Fitch
2	William Diaz
3	Kevin Rodriguez
4	Catherine Crandell
5	Richard Tillman

Kafka topics: Flink ingestion

- Data acquisition in Flink can be automated
- Design the topics ingestion in Flink:
 - Flink streaming tables
 - queried with FlinkSQL continuous queries, recorded in Flink
 - Kafka connector for Flink ([Table & SQL API](#)):
 - files:
 - ***sql-client-conf.yaml***: Kafka → Flink
 - ***sql-gateway-defaults.yaml***: Flink JDBC Gateway
 - **table schema**: topics fields, datatypes, watermarks, ...
 - **connector properties**: Kafka address, schema registry, ...

Example: DEALER2_VEHICLES topic

- name: D2_VEHICLES ← **TABLE NAME**

type: source

update-mode: append

data: SOURCE → FLINK

schema:

- name: userID

type: BIGINT

- name: rid

type: BIGINT

- name: type

type: STRING

- name: manufacturer

type: STRING

- name: model

type: STRING

- name: plate

type: STRING

- name: status

type: STRING

- name: ts


type: STRING

append new data, when is available from the source (Kafka)

TABLE SCHEMA

Reference: [Apache Kafka SQL Connector](#)

Example: DEALER2_VEHICLES topic



Relational Streaming Data Integration...

Now, we have a Flink streaming table for each Kafka topic

- *DEALER1_CARS* and *DEALER1_TRUCKS*
- *DEALER2_VEHICLES* and *DEALER2_USERS*

The data streams are still not integrated!!!

Relational Streaming Data Integration...

Now, we have a Flink streaming table for each Kafka topic

- *DEALER1_CARS* and *DEALER1_TRUCKS*
- *DEALER2_VEHICLES* and *DEALER2_USERS*

We can use **OntopStream** to create a **unified logical view** of the data streams...

Flink relational streams:

- exposed to OntopStream using the ***Flink JDBC Gateway***
- can be queried with ***FlinkSQL*** continuous queries

Relational Streaming Data Integration...

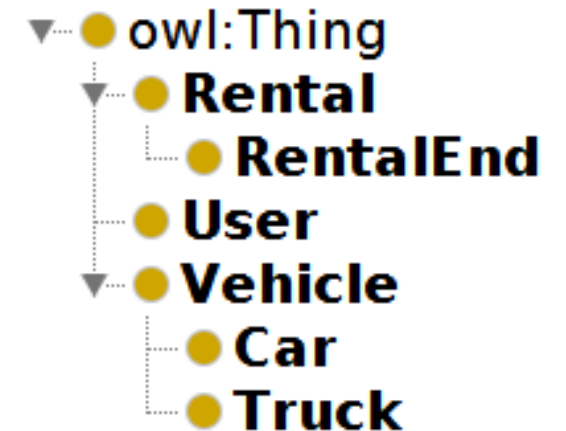
- Onpstream automates:
 - **RSP-QL → FlinkSQL** query rewriting
 - **relational → RDF** response streams translation
- To use OntopStream for the streaming data integration tasks we need:
 1. **Ontology**: provides the unified logical view to the user
 - Classes
 - Object Properties
 - Data Properties
 2. **Streaming-VKG mappings**: bridges the ontology with data streams (Kafka messages in Flink)
 3. **JDBC connection** configuration

1) Ontology Design

Ontology design made using [Protégé](#), an open-source graphical ontology editor.

Classes:


- express the **logical concepts** of the unified logical view
- The **Car** and **Truck** concepts are expressed as subclasses of **Vehicle**
i.e., a *Tesla Model X* is a Car, but also a transportation Vehicle
- **RentalEnd** is a specialization of (subclass) **Rental**
it will be useful later for queries about ended rentals



1) Ontology Design


Object Properties:

- ease the mapping process
- express implicit domain/range restrictions on Class instances:
 - the **user** property range is **User**
 - the **vehicle** property range is **Vehicle**



Data Properties:

- expose the Kafka messages entries
 - Vehicle details (manufacturer, model)
 - Timestamps
 - Users personal information (name)



2) Streaming-VKG mappings

OntopStream answers **RSP-QL queries** with **RDF streams** of semantically-enriched responses based on:

- ontological concepts
- relational data streams: retrieved through **Streaming VKG queries** registered in Flink

Streaming-VKG mapping

- binding between a set of RDF statements and FlinkSQL selection query
- connects the ontological layer terms to data streams (in this tutorial, Kafka messages in Flink)
- consists of:
 - **MappingID**: friendly name to identify the mapping
 - **Source**: FlinkSQL query for the data extraction from the Flink streaming tables
 - **Target**: one or more RDF statements corresponding to the VKG generated by the single entry obtained from the data extracted with the Source query

2) Streaming-VKG mappings: Baranch A

Entities:

- **Rental:** each rental ID in the stream
- **Vehicle:** plate numbers
- **User:** client names

Kind of rented vehicle?

- **D1_CARS** table stores **Cars** data
- **D1_TRUCKS** table stores **Trucks** data

Start or ended lease?

- the **status** field refers to the rental state
- we can use a **WHERE** clause in the source query to filter out rentals by their status:
 - **status='START'** retrieves the starting rentals Kafka messages
 - **status='END'** retrieves the ended rentals Kafka messages

2) Streaming-VKG mappings: Baranch A

RentalEnd subclass of Rental:

- to ease the complexity of queries asking only for ended rentals, we use the subclass specialization

Started rentals [rentals.obda]

mappingId	DEALER1-CarRental
target	:D1_C{rid} a :Rental; :user :{user}; :hasStart {ts}^^xsd:dateTime; :car :{plate}. :{plate} a :Car; :manufacturer {manufacturer}; :model {model}.
source	SELECT rid, user, ts, plate, manufacturer, model FROM D1_CARS WHERE status='START'
mappingId	DEALER1-TruckRental
target	:D1_T{rid} a :Rental; :user :{user}; :hasStart {ts}^^xsd:dateTime; :truck :{plate}. :{plate} a :Truck; :manufacturer {manufacturer}; :model {model}.
source	SELECT rid, user, ts, plate, manufacturer, model FROM D1_TRUCKS WHERE status='START'

Ended rentals [rentals.obda]

mappingId	DEALER1-CarRentalEnd
target	:D1_C{rid} a :RentalEnd; :hasEnd {ts}^^xsd:dateTime; :car :{plate}.
source	SELECT rid,ts,plate FROM D1_CARS WHERE status='END'
mappingId	DEALER1-TruckRentalEnd
target	:D1_T{rid} a :RentalEnd; :hasEnd {ts}^^xsd:dateTime; :truck :{plate}.
source	SELECT rid,ts,plate FROM D1_TRUCKS WHERE status='END'

2) Streaming-VKG mappings: Branch B

Entities:

- **Rental:** each rental ID in the stream
- **Vehicle:** plate numbers
- **User:** client names

Kind of rented vehicle?

- the **type** field refers to the kind of vehicle in the **D2_VEHICLES** table
- for starting rentals, we can use a **WHERE** clause in the source query to determine the vehicle:
 - **type= 'Car'** retrieves **Car** rental entries
 - **type= 'Truck'** retrieves **Truck** rental entries
- for ending rentals, since the vehicle class is determined in the starting rental messages:
 - use the generic **vehicle** object property (property range is **Vehicle**)

Start or ended lease? (same as Branch A)

- use the **WHERE** clause in the source query to filter out rentals by their **status** field

2) Streaming-VKG mappings: Branch B

Users are kept in a separate topic:

- need to combine the Flink streaming tables **D2_VEHICLES** and **D2_USERS**
- FlinkSQL source query with a **JOIN** over the **userID** field

RentalEnd subclass of **Rental**: (same as Branch A)

- to ease the complexity of queries asking only for ended rentals, we use the subclass specialization

Started/Ended rentals [rentals.obda]

mappingId	DEALER2-CarRental
target	:D2_{rid} a :Rental; :user {name}; :hasStart {ts}^^xsd:dateTime; :car {plate}. {plate} a :Car; :manufacturer {manufacturer}; :model {model}.
source	SELECT rid,name,ts,plate,manufacturer,model FROM D2_VEHICLES,D2_USERS WHERE D2_VEHICLES.userID=D2_USERS.userID AND type='Car' AND status='START'
mappingId	DEALER2-TruckRental
target	:D2_{rid} a :Rental; :user {name}; :hasStart {ts}^^xsd:dateTime; :truck {plate}. {plate} a :Truck; :manufacturer {manufacturer}; :model {model}.
source	SELECT rid,name,ts,plate,manufacturer,model FROM D2_VEHICLES,D2_USERS WHERE D2_VEHICLES.userID=D2_USERS.userID AND type='Truck' AND status='START'
mappingId	DEALER2-RentalEnd
target	:D2_{rid} a :RentalEnd; :hasEnd {ts}^^xsd:dateTime; :vehicle {plate}.
source	SELECT rid,ts,plate FROM D2_VEHICLES,D2_USERS WHERE D2_VEHICLES.userID=D2_USERS.userID AND status='END'

3) JDBC connection

- OntopStream interacts with Apache Flink:
 - through **JDBC calls**
 - using a **custom JDBC driver**
- Before starting OntopStream, we need to configure the connection to the **Flink JDBC Gateway**
- The configuration must be specified in a **property file**, passed as input to OntopStream on its startup

rentals.property

```
jdbc.url=jdbc:flink://sql-client:8083?planner=blink
jdbc.driver=com.ververica.flink.table.jdbc.FlinkDriver
jdbc.user=
jdbc.name=test-RSE-streaming
jdbc.fetchSize=1
jdbc.password=
```

Tutorial: OntopStream hands-on

OntopStream startup

- The OntopStream docker image is available on DockerHub

hub.docker.com/r/chimerasuite/ontop-stream

- We can now start the OntopStream endpoint using the command:

```
docker-compose -f ontop.yml up -d
```

- If we look at the configuration in the `ontop.yml` file we can see the three input files:
 - `tutorial.owl`: contained the ontology describing the user unified logical view
 - `tutorial.obda`: the Streaming-VKG mappings we've designed
 - `tutorial.properties`: the JDBC connection properties

JupyterLab setup

- We're finally ready for querying the streams of data using a python notebook
- The platform is accessible from `http://<IP-ADDRESS>:8888/lab?token=TEST`

The screenshot shows the JupyterLab web interface. On the left is a file explorer sidebar with a search bar and a table of files. The table has two columns: 'Name' and 'Last Modified'. The files listed are 'data_generator', 'output', 'SPARQLStreamWrapper', 'dealer_1.ipynb', 'dealer_2.ipynb', and 'OntopStream_tutorial.ipynb', all modified '5 days ago'. 'dealer_1.ipynb' is selected. On the right is a notebook editor for 'dealer_1.ipynb' with the title 'Kafka Producer: dealer #1'. The notebook contains two code cells. The first cell has the command `pip install confluent_kafka --quiet`. The second cell contains Python code for creating Kafka topics. Arrows from the text labels on the left point to the corresponding elements in the interface.

Kafka producer #1 points to the `dealer_1.ipynb` file in the file explorer.

Kafka producer #2 points to the `dealer_2.ipynb` file in the file explorer.

OntopStream demo points to the `OntopStream_tutorial.ipynb` file in the file explorer.

```
[ ]: pip install confluent_kafka --quiet

[ ]: from confluent_kafka import Producer, KafkaError
      from confluent_kafka.admin import AdminClient, NewTopic, NewPartitions
      import string

      topic_cars = "DEALER1_CARS"
      topic_trucks = "DEALER1_TRUCKS"

      a = AdminClient({'bootstrap.servers': 'kafka:9092'})
      fs = a.create_topics([NewTopic(topic_cars, num_partitions=1, replication_factor=1),
                           NewTopic(topic_trucks, num_partitions=1, replication_factor=1)])

      for t, f in fs.items():
          try:
              f.result() # The result itself is None
              print("Topic {} created".format(t))
          except Exception as e:
              print("Failed to create topic {}: {}".format(t, e))
```

First Query

Get the car rentals (from both the branches)

```
from SPARQLStreamWrapper import SPARQLStreamWrapper, TSV

sparql = SPARQLStreamWrapper("http://ontop:8080/sparql")
sparql.setQuery("""
PREFIX : <http://www.semanticweb.org/car-rental#>

SELECT ?user ?car ?model ?start

WHERE {
    ?car a :Car; :model ?model.
    ?rent a :Rental; :car ?car.
    ?rent :hasStart ?start; :user ?user.
}
""")

sparql.addParameter("streaming-mode", "single-element")
sparql.setReturnFormat(TSV)

results=sparql.query()

try:
    for result in results:
        data = result.getRawResponse().decode('utf8')
        print(data)
except KeyboardInterrupt:
    sparql.endQuery()
    print("Ended by user")
```

?user ?car ?model ?start

<http://www.semanticweb.org/car-rental#Molly%20Davis> <http://www.semanticweb.org/car-rental#FJ7PUJJ> Panda "2022-03-31 09:52:30"^^<http://www.w3.org/2001/XMLSchema#dateTime>

<http://www.semanticweb.org/car-rental#Laura%20Baker> <http://www.semanticweb.org/car-rental#JFGJ60A> Model S "2022-03-31 09:52:54"^^<http://www.w3.org/2001/XMLSchema#dateTime>

<http://www.semanticweb.org/car-rental#Kevin%20Rodriguez> <http://www.semanticweb.org/car-rental#DFU4HJF> A3 "2022-03-31 09:52:35"^^<http://www.w3.org/2001/XMLSchema#dateTime>

<http://www.semanticweb.org/car-rental#Catherine%20Crandell> <http://www.semanticweb.org/car-rental#784JD93> Classe C "2022-03-31 09:53:13"^^<http://www.w3.org/2001/XMLSchema#dateTime>

First Query

Get the car rentals (from both the branches)

```
from SPARQLStreamWrapper import SPARQLStreamWrapper, TSV

sparql = SPARQLStreamWrapper("http://ontop:8080/sparql")
sparql.setQuery("""
PREFIX : <http://www.semanticweb.org/car-rental#>

SELECT ?user ?car ?model ?start

WHERE {
    ?car a :Car; :model ?model.
    ?rent a :Rental; :car ?car.
    ?rent :hasStart ?start; :user ?user.
}
""")

sparql.addParameter("streaming-mode", "single-element")
sparql.setReturnFormat(TSV)

results=sparql.query()

try:
    for result in results:
        data = result.getRawResponse().decode('utf8')
        print(data)
except KeyboardInterrupt:
    sparql.endQuery()
    print("Ended by user")
```

```
?user    ?car    ?model  ?start

<http://www.semanticweb.org/car-rental#Molly%20Davis> <http://www.semanticweb.org/car-rental#FJ7PUJJ> Panda    "2022-03-31 09:52:30"^^<http://www.w3.org/2001/XMLSchema#dateTime>
<http://www.semanticweb.org/car-rental#Laura%20Baker> <http://www.semanticweb.org/car-rental#JFGJ60A> Model S "2022-03-31 09:52:54"^^<http://www.w3.org/2001/XMLSchema#dateTime>
<http://www.semanticweb.org/car-rental#Kevin%20Rodriguez> <http://www.semanticweb.org/car-rental#DFU4HJF> A3      "2022-03-31 09:52:35"^^<http://www.w3.org/2001/XMLSchema#dateTime>
<http://www.semanticweb.org/car-rental#Catherine%20Crandell> <http://www.semanticweb.org/car-rental#784JD93> Classe C    "2022-03-31 09:53:13"^^<http://www.w3.org/2001/XMLSchema#dateTime>
<http://www.semanticweb.org/car-rental#William%20Diaz> <http://www.semanticweb.org/car-rental#FGL1X62> Tipo      "2022-03-31 09:53:33"^^<http://www.w3.org/2001/XMLSchema#dateTime>
<http://www.semanticweb.org/car-rental#Douglas%20Fitch> <http://www.semanticweb.org/car-rental#UF94JF> "911"    "2022-03-31 09:53:53"^^<http://www.w3.org/2001/XMLSchema#dateTime>
<http://www.semanticweb.org/car-rental#William%20Diaz> <http://www.semanticweb.org/car-rental#AL3SLS> A4       "2022-03-31 09:54:25"^^<http://www.w3.org/2001/XMLSchema#dateTime>
```

Second Query – Real-Time Filtering

Get the Porsche and Tesla car rentals (from both the branches)

```
from SPARQLStreamWrapper import SPARQLStreamWrapper, TSV

sparql = SPARQLStreamWrapper("http://ontop:8080/sparql")
sparql.setQuery("""
PREFIX : <http://www.semanticweb.org/car-rental#>

SELECT ?user ?car ?man ?model ?start

WHERE {
  ?car a :Car; :model ?model; :manufacturer ?man.
  ?rent a :Rental; :car ?car.
  ?rent :hasStart ?start; :user ?user.
  FILTER(?man="Tesla" || ?man="Porsche")
}
""")

sparql.addParameter("streaming-mode", "single-element")
sparql.setReturnFormat(TSV)

results=sparql.query()

try:
    for result in results:
        data = result.getRawResponse().decode('utf8')    # Get response from OntopStream
        data = data.replace("%20", " ")                  # Clean IDs
        print(data)
except KeyboardInterrupt:
    sparql.endQuery()
    print("Ended by user")
```

real-time filtering condition, translated in a WHERE clause over the queried Flink Dynamic Tables

```
?user    ?car    ?man    ?model  ?start

<http://www.semanticweb.org/car-rental#Laura Baker> <http://www.semanticweb.org/car-rental#JFGJ60A> Tesla Model S "2022-03-31 09:52:54"^^<http://www.w3.org/2001/XMLSchema#dateTime>
<http://www.semanticweb.org/car-rental#Frank Cover> <http://www.semanticweb.org/car-rental#JFGJ60A> Tesla Model S "2022-03-31 10:02:57"^^<http://www.w3.org/2001/XMLSchema#dateTime>
<http://www.semanticweb.org/car-rental#Douglas Fitch> <http://www.semanticweb.org/car-rental#DR7TGF0> Tesla Model X "2022-03-31 09:58:18"^^<http://www.w3.org/2001/XMLSchema#dateTime>
```

Second Query – Real-Time Filtering

Get the Porsche and Tesla car rentals (from both the branches)

```
from SPARQLStreamWrapper import SPARQLStreamWrapper, TSV

sparql = SPARQLStreamWrapper("http://ontop:8080/sparql")
sparql.setQuery("""
PREFIX : <http://www.semanticweb.org/car-rental#>

SELECT ?user ?car ?man ?model ?start

WHERE {
  ?car a :Car; :model ?model; :manufacturer ?man.
  ?rent a :Rental; :car ?car.
  ?rent :hasStart ?start; :user ?user.
  FILTER(?man="Tesla" || ?man="Porsche")
}
""")

sparql.addParameter("streaming-mode", "single-element")
sparql.setReturnFormat(TSV)

results=sparql.query()

try:
    for result in results:
        data = result.getRawResponse().decode('utf8')    # Get response from OntopStream
        data = data.replace("%20", " ")                  # Clean IDs
        print(data)
except KeyboardInterrupt:
    sparql.endQuery()
    print("Ended by user")
```

real-time filtering condition, translated in a WHERE clause over the queried Flink Dynamic Tables

?user	?car	?man	?model	?start
<http://www.semanticweb.org/car-rental#Laura Baker>	<http://www.semanticweb.org/car-rental#JFGJ60A>	Tesla	Model S	"2022-03-31 09:52:54"^^<http://www.w3.org/2001/XMLSchema#dateTime>
<http://www.semanticweb.org/car-rental#Frank Cover>	<http://www.semanticweb.org/car-rental#JFGJ60A>	Tesla	Model S	"2022-03-31 10:02:57"^^<http://www.w3.org/2001/XMLSchema#dateTime>
<http://www.semanticweb.org/car-rental#Douglas Fitch>	<http://www.semanticweb.org/car-rental#DR7TGF0>	Tesla	Model X	"2022-03-31 09:58:18"^^<http://www.w3.org/2001/XMLSchema#dateTime>
<http://www.semanticweb.org/car-rental#Lucille Bouchard>	<http://www.semanticweb.org/car-rental#8NMSMII>	Tesla	Model X	"2022-03-31 10:00:59"^^<http://www.w3.org/2001/XMLSchema#dateTime>
<http://www.semanticweb.org/car-rental#Douglas Fitch>	<http://www.semanticweb.org/car-rental#AB7TGX0>	Tesla	Model Y	"2022-03-31 10:01:27"^^<http://www.w3.org/2001/XMLSchema#dateTime>
<http://www.semanticweb.org/car-rental#Douglas Fitch>	<http://www.semanticweb.org/car-rental#UF94JF>	Porsche	"911"	"2022-03-31 09:53:53"^^<http://www.w3.org/2001/XMLSchema#dateTime>

Third Query – Reasoning...

Get the rentals for Mercedes vehicles (trucks and cars), persist the results in a CSV file

```
from SPARQLStreamWrapper import SPARQLStreamWrapper, CSV
import os
```

```
sparql = SPARQLStreamWrapper("http://ontop:8080/sparql")
sparql.setQuery("""
PREFIX : <http://www.semanticweb.org/car-rental#>
```

```
SELECT ?user ?plate ?model ?start
```

```
WHERE {
  ?plate a :Vehicle; :manufacturer ?man; :model ?model.
  ?rent a :Rental; :vehicle ?plate.
  ?rent :hasStart ?start; :user ?user.
  FILTER(?man="Mercedes")
}
```

```
""")
sparql.addParameter("streaming-mode", "single-element")
sparql.setReturnFormat(CSV)
```

```
file=open("output/query_3.csv", "w+")
```

```
results=sparql.query()
```

```
try:
```

```
    for result in results:
```

```
        data = result.getRawResponse().decode('utf8')
```

```
        data = data.replace("http://www.semanticweb.org/car-rental#", "")
```

```
        data = data.replace("%20", " ")
```

```
        print(data)
```

```
        file.write(data)
```

```
        file.flush()
```

```
        os.fsync(file.fileno())
```


```
except KeyboardInterrupt:
```

```
    sparql.endQuery()
```

```
    file.close()
```

```
    print("Ended by user")
```

Real-time reasoning, based on the mapping ontology. *Cars and Trucks are both vehicles*



Query results persistence

```
# Get response from OntopStream
```

```
# Remove prefixes
```

```
# Clean Names
```

```
# Write response in the file
```

```
# Flush the writing operation
```

Third Query – Reasoning...

user,plate,model,start

Catherine Crandell,784JD93,Classe C,2022-03-31 11:01:16

Kevin Rodriguez,KD94KDS,Vito,2022-03-31 11:04:49

Wayne Flower,012JKD0,Vito,2022-03-31 11:05:10

Laura Baker,B38SDJA,Citan,2022-03-31 11:06:11

Mark Haws,D74HJDK,Classe E,2022-03-31 11:08:43

Jeanie Morgan,012JKD0,Vito,2022-03-31 11:09:40

Kevin Rodriguez,B38SDJA,Citan,2022-03-31 11:12:21

Catherine Crandell,KD94KDS,Vito,2022-03-31 11:18:08

Car (branch 1)

Truck (branch 2)

query_3.csv				
Delimiter: ,				
	user	plate	model	start
1	Catherine Crandell	784JD93	Classe C	2022-03-31 11:01:16
2	Kevin Rodriguez	KD94KDS	Vito	2022-03-31 11:04:49
3	Wayne Flower	012JKD0	Vito	2022-03-31 11:05:10
4	Laura Baker	B38SDJA	Citan	2022-03-31 11:06:11
5	Mark Haws	D74HJDK	Classe E	2022-03-31 11:08:43
6	Jeanie Morgan	012JKD0	Vito	2022-03-31 11:09:40
7	Kevin Rodriguez	B38SDJA	Citan	2022-03-31 11:12:21
8	Catherine Crandell	KD94KDS	Vito	2022-03-31 11:18:08

Fourth Query – Windowing (future developments)

Get the trucks old rentals (rentals which have been finished)

```
from SPARQLStreamWrapper import SPARQLStreamWrapper, CSV
import os

sparql = SPARQLStreamWrapper("http://ontop:8080/sparql")
sparql.setQuery("""
PREFIX : <http://www.semanticweb.org/car-rental#>

SELECT ?rent ?manuf ?model ?end
FROM NAMED WINDOW :wind1 ON :trips [RANGE PT1M STEP PT1M]
WHERE {
    ?truck a :Truck; :manufacturer ?manuf; :model ?model.
    ?rent a :RentalEnd; :truck ?truck.
    ?rent :hasEnd ?end.
}
""")

sparql.addParameter("streaming-mode", "single-element")
sparql.setReturnFormat(CSV)

file=open("output/query_4.csv", "w+")

results=sparql.query()

try:
    for result in results:
        data = result.getRawResponse().decode('utf8')
        data = data.replace("http://www.semanticweb.org/car-rental#", "")
        print(data)
        file.write(data)
        file.flush()
        os.fsync(file.fileno())
except KeyboardInterrupt:
    sparql.endQuery()
    file.close()
    print("Ended by user")
```

RSP-QL window condition...

Responses @t1

```
rent,manuf,model,end
D1_T5,Fiat,Ducato,2022-03-31 10:47:55
D2_8,Mercedes,Vito,2022-03-31 10:48:12
D1_T2,Fiat,Ducato,2022-03-31 10:49:40
D1_T1,Iveco,Daily,2022-03-31 10:50:01
D2_13,Mercedes,Vito,2022-03-31 10:50:02
D2_9,Mercedes,Citan,2022-03-31 10:50:30
```


Fourth Query – Windowing (future developments)

Get the trucks old rentals (rentals which have been finished)

```
from SPARQLStreamWrapper import SPARQLStreamWrapper, CSV
import os

sparql = SPARQLStreamWrapper("http://ontop:8080/sparql")
sparql.setQuery("""
PREFIX : <http://www.semanticweb.org/car-rental#>

SELECT ?rent ?manuf ?model ?end
FROM NAMED WINDOW :wind1 ON :trips [RANGE PT1M STEP PT1M]
WHERE {
    ?truck a :Truck; :manufacturer ?manuf; :model ?model.
    ?rent a :RentalEnd; :truck ?truck.
    ?rent :hasEnd ?end.
}
""")

sparql.addParameter("streaming-mode", "single-element")
sparql.setReturnFormat(CSV)

file=open("output/query_4.csv", "w+")

results=sparql.query()

try:
    for result in results:
        data = result.getRawResponse().decode('utf8')
        data = data.replace("http://www.semanticweb.org/car-rental#", "")
        print(data)
        file.write(data)
        file.flush()
        os.fsync(file.fileno())
except KeyboardInterrupt:
    sparql.endQuery()
    file.close()
    print("Ended by user")
```

RSP-QL window condition...

Responses @t2

```
rent,manuf,model,end
D1_T5,Fiat,Ducato,2022-03-31 10:47:55
D2_8,Mercedes,Vito,2022-03-31 10:48:12
D1_T2,Fiat,Ducato,2022-03-31 10:49:40
D1_T1,Iveco,Daily,2022-03-31 10:50:01
D2_13,Mercedes,Vito,2022-03-31 10:50:02
D2_9,Mercedes,Citan,2022-03-31 10:50:30
D1_T3,Iveco,Daily,2022-03-31 10:51:57
D2_7,Mercedes,Vito,2022-03-31 10:52:20
```

Thank you !!